

# Exceptional Control Flow: Exceptions and Processes

Read Chap 8.1-8.4

Instructor: Jennifer Wong-Ma  
jwongma@uci.edu

# Today


- **Exceptional Control Flow**
- Exceptions
- Processes
- Process Control

# Control Flow

## ■ Processors do only one thing:

- From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
- This sequence is the CPU's *control flow* (or *flow of control*)

### *Physical control flow*



# Altering the Control Flow

## ■ Up to now: two mechanisms for changing control flow:

- Jumps and branches
- Call and return

Both react to changes in *program state*

## ■ Insufficient for a useful system:

Difficult to react to changes in *system state*

- Data arrives from a disk or a network adapter
- Instruction divides by zero
- User hits Ctrl-C at the keyboard
- System timer expires

## ■ System needs mechanisms for “exceptional control flow”

# Exceptional Control Flow

- Exists at all levels of a computer system

- Low level mechanisms

- 1. **Exceptions**
  - Change in control flow in response to a system event (i.e., change in system state)
  - Implemented using combination of hardware and OS software

- Higher level mechanisms

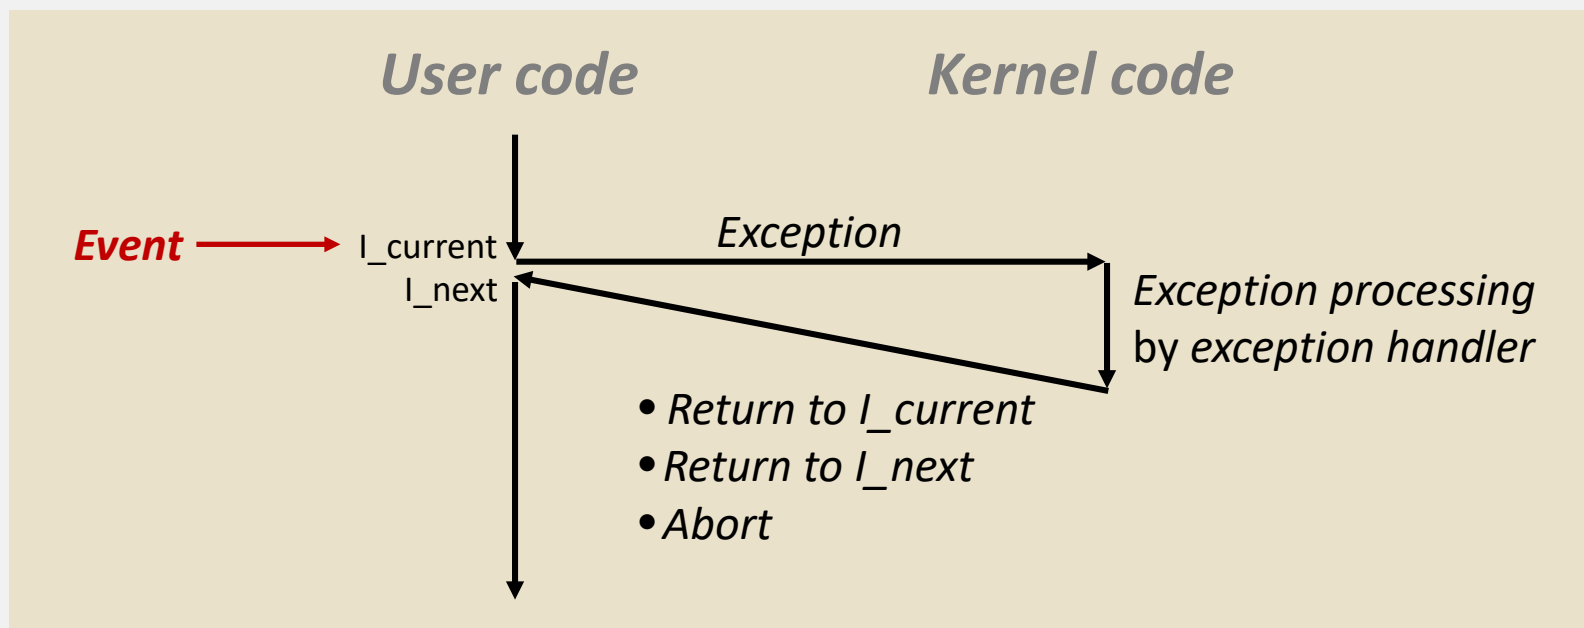
- 2. **Process context switch**
  - Implemented by OS software and hardware timer
- 3. **Signals**
  - Implemented by OS software
- 4. **Nonlocal jumps**: `setjmp( )` and `longjmp( )`
  - Implemented by C runtime library

# Today

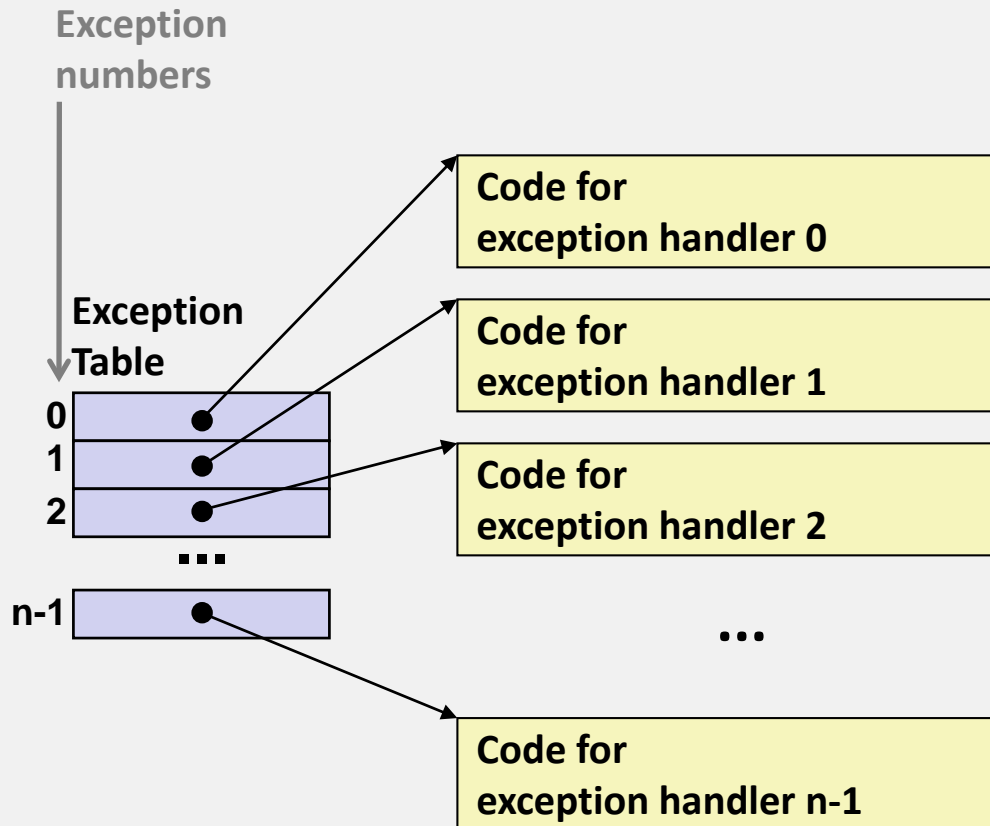
- Exceptional Control Flow
- **Exceptions**
- Processes
- Process Control

# Exceptions

- An **exception** is a transfer of control to the OS *kernel* in response to some *event* (i.e., change in processor state)
  - Kernel is the memory-resident part of the OS
  - Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C



# Exception Tables



- Each type of event has a unique exception number  $k$
- $k$  = index into exception table (a.k.a. interrupt vector)
- Handler  $k$  is called each time exception  $k$  occurs



# Asynchronous Exceptions (Interrupts)

## ■ Caused by events external to the processor

- Indicated by setting the processor's *interrupt pin*
- Handler returns to “next” instruction

## ■ Examples:

- Timer interrupt
  - Every few ms, an external timer chip triggers an interrupt
  - Used by the kernel to take back control from user programs
- I/O interrupt from external device
  - Hitting Ctrl-C at the keyboard
  - Arrival of a packet from a network
  - Arrival of data from a disk

# Synchronous Exceptions

## ■ Caused by events that occur as a result of executing an instruction:

- **Traps**

- Intentional
- Examples: **system calls**, breakpoint traps, special instructions
- Returns control to “next” instruction

- **Faults**

- Unintentional but possibly recoverable
- Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
- Either re-executes faulting (“current”) instruction or aborts

- **Aborts**

- Unintentional and unrecoverable
- Examples: illegal instruction, parity error, machine check
- Aborts current program

# Examples of x86-64 Exceptions

<i>Exception Number</i>	<i>Description</i>	<i>Exception Class</i>
0	Divide by zero	Fault
13	General protection fault	Fault
14	Page fault	Fault
18	Machine check	Abort
32-255	OS-defined exceptions	Interrupt or trap

# System Calls

- Each x86-64 system call has a unique ID number
- Examples:

<i>Number</i>	<i>Name</i>	<i>Description</i>
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

# System Call Example: Opening File

- **User calls:** `open(filename, options)`
- **Calls `__open` function, which invokes system call instruction `syscall`**

```
0000000000e5d70 <__open>:
```

```
...
```

```
e5d79: b8 02 00 00 00    mov $0x2,%eax # open is syscall #2
```

```
e5d7e: 0f 05            syscall      # Return value in %rax
```


```
e5d80: 48 3d 01 f0 ff ff  cmp $0xfffffffffff001,%rax
```

```
...
```

```
e5dfa: c3              retq
```

*User code*

*Kernel code*



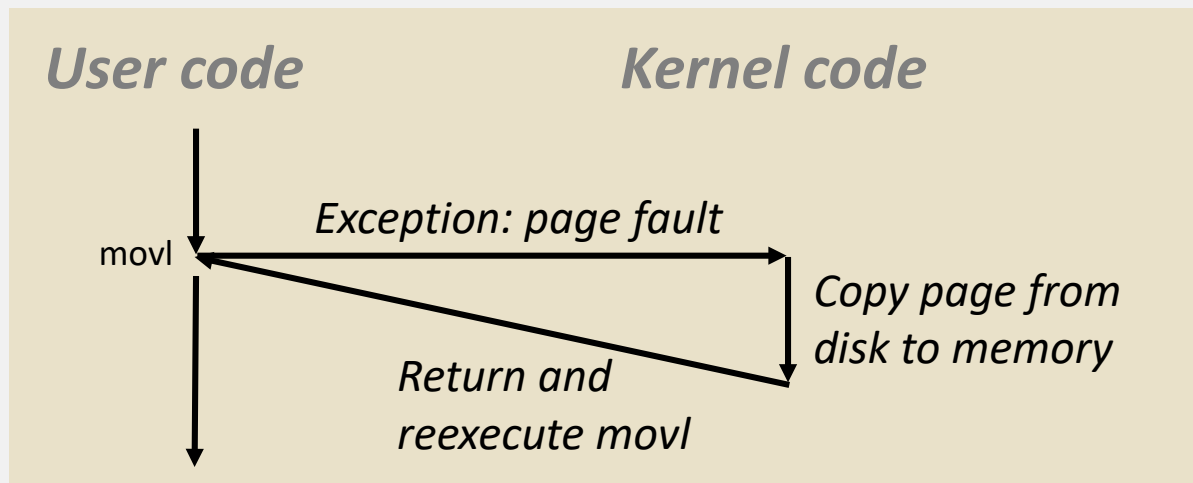
- `%rax` contains syscall number
- Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
- Return value in `%rax`
- Negative value is an error corresponding to negative `errno`

# Fault Example: Page Fault

- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```


80483b7:	c7 05 10 9d 04 08 0d	movl	\$0xd,0x8049d10
----------	----------------------	------	-----------------



# Fault Example: Invalid Memory Reference

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```

80483b7: c7 05 60 e3 04 08 0d movl \$0xd,0x804e360



- Sends SIGSEGV signal to user process
- User process exits with “segmentation fault”

# Today

- Exceptional Control Flow
- Exceptions
- **Processes**
- Process Control




# Processes

- **Definition:** A *process* is an instance of a running program.

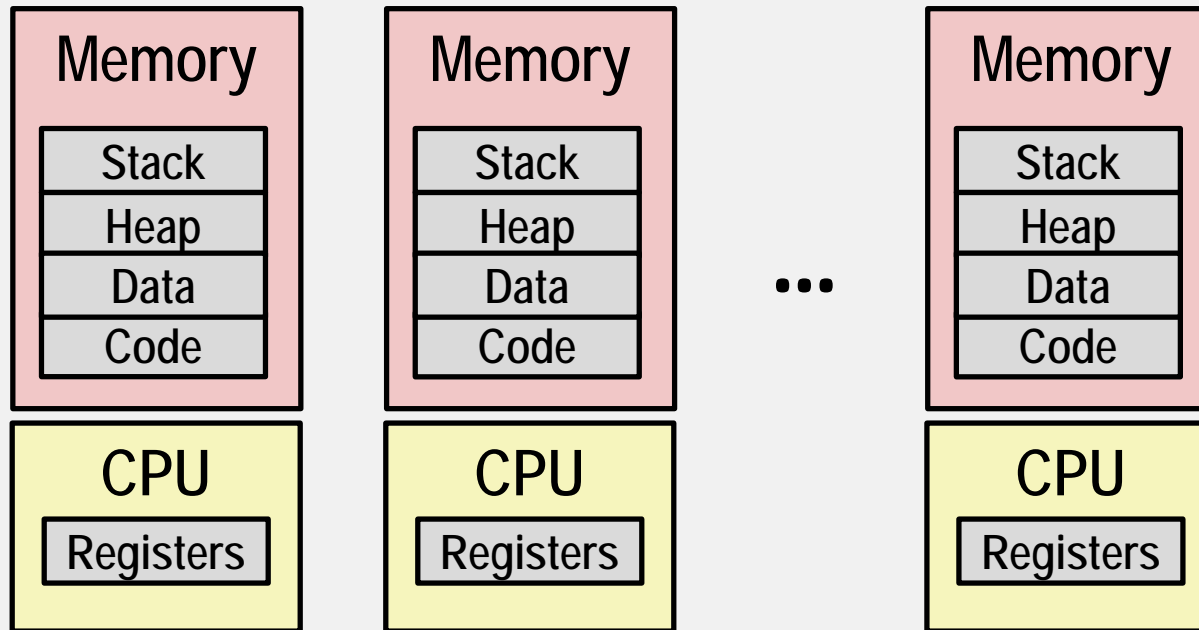
- One of the most profound ideas in computer science
- Not the same as “program” or “processor”

- **Process provides each program with two key abstractions:**

- *Logical control flow*
  - Each program seems to have exclusive use of the CPU
  - Provided by kernel mechanism called *context switching*
- *Private address space*
  - Each program seems to have exclusive use of main memory.
  - Provided by kernel mechanism called *virtual memory*



# Multiprocessing: The Illusion



## ■ Computer runs many processes simultaneously

- Applications for one or more users
  - Web browsers, email clients, editors, ...
- Background tasks
  - Monitoring network & I/O devices

# Multiprocessing Example

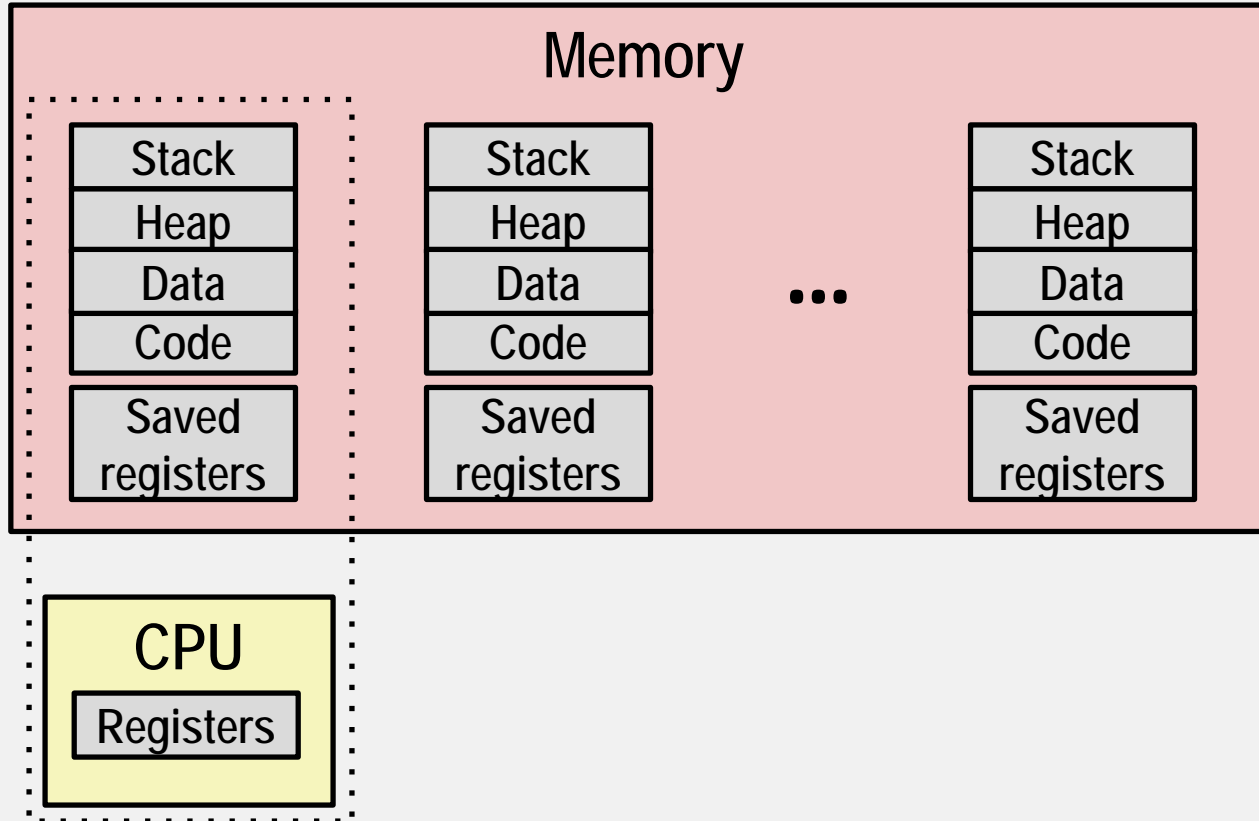
```
Processes: 123 total, 5 running, 9 stuck, 109 sleeping, 611 threads
Load Avg: 1.03, 1.13, 1.14 CPU usage: 3.27% user, 5.15% sys, 91.56% idle
SharedLibs: 576K resident, 0B data, 0B linkedit.
MemRegions: 27958 total, 1127M resident, 35M private, 494M shared.
PhysMem: 1039M wired, 1974M active, 1062M inactive, 4076M used, 18M free.
VM: 280G vsize, 1091M framework vsize, 23075213(1) pageins, 5843367(0) pageouts.
Networks: packets: 41046228/11G in, 66083096/77G out.
Disks: 17874391/349G read, 12847373/594G written.

PID    COMMAND    %CPU TIME    #TH    #WQ    #PORT    #MREG    RPRVT    RSHRD    RSIZE    VPRVT    VSIZE
99217-  Microsoft Of 0.0 02:28.34 4      1      202     418     21M     24M     21M     66M     763M
99051  usbmuxd     0.0 00:04.10 3      1      47      66     436K    216K    480K    60M     2422M
99006  iTunesHelper 0.0 00:01.23 2      1      55      78     728K    3124K   1124K   43M     2429M
84286  bash        0.0 00:00.11 1      0      20      24     224K    732K    484K    17M     2378M
84285  xterm       0.0 00:00.83 1      0      32      73     656K    872K    692K    9728K   2382M
55939-  Microsoft Ex 0.3 21:58.97 10     3      360     954     16M     65M     46M     114M    1057M
54751  sleep       0.0 00:00.00 1      0      17      20     92K     212K    360K    9632K   2370M
54739  launchdadd  0.0 00:00.00 2      1      33      50     488K    220K    1736K   48M     2409M
54737  top         6.5 00:02.53 1/1    0      30      29     1416K   216K    2124K   17M     2378M
54719  automountd  0.0 00:00.02 7      1      53      64     860K    216K    2184K   53M     2413M
54701  ocspd       0.0 00:00.05 4      1      61      54     1268K   2644K   3132K   50M     2426M
54661  Grab        0.6 00:02.75 6      3      222+    389+    15M+    26M+    40M+    75M+    2556M+
54659  cookiec     0.0 00:00.15 2      1      40      61     3316K   224K    4088K   42M     2411M
53818  mdworker    0.0 00:01.67 4      1      52      91     7628K   7412K   16M     48M     2438M
50878  mdworker    0.0 00:11.17 3      1      53      91     2464K   6148K   9976K   44M     2434M
50410  xterm       0.0 00:00.13 1      0      32      73     280K    872K    532K    9700K   2382M
50078  emacs       0.0 00:06.70 1      0      20      35     52K     216K    88K     18M     2392M
```

## ■ Running program “top” on Mac

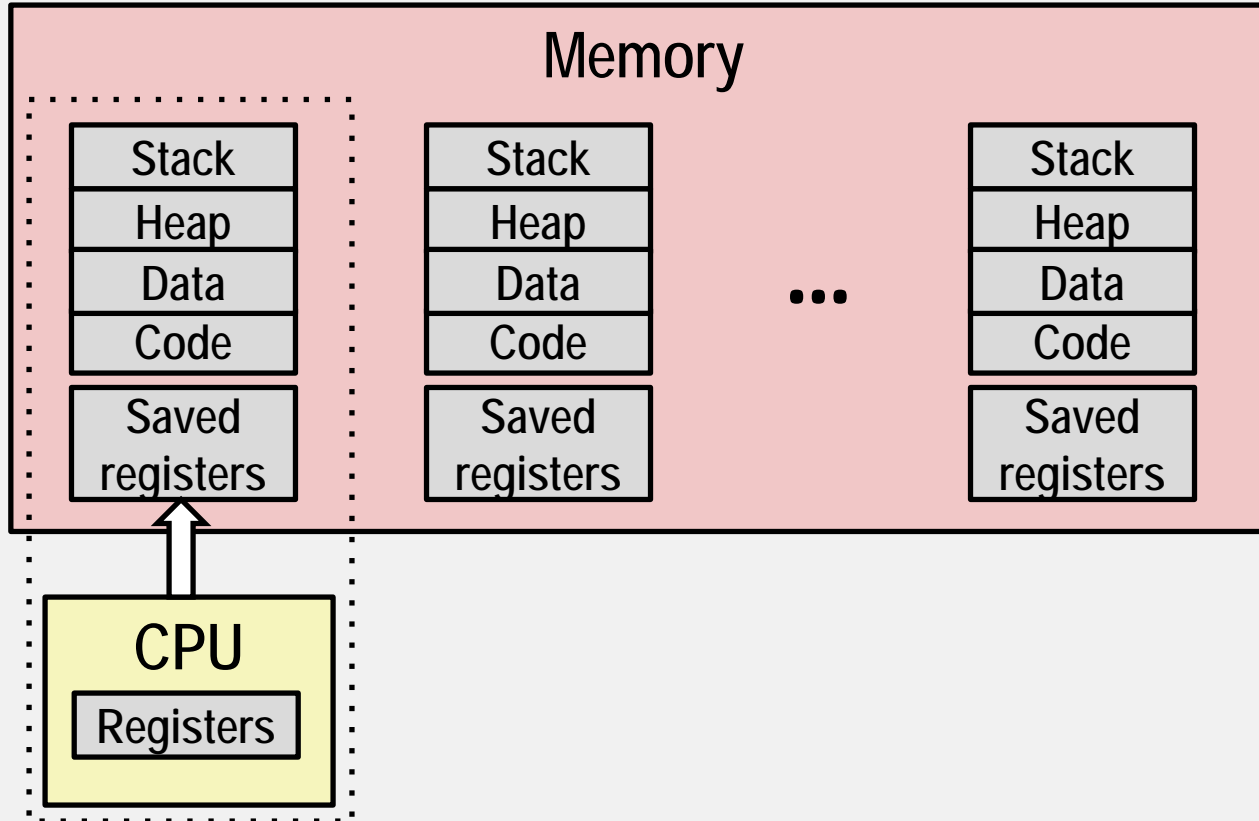
- System has 123 processes, 5 of which are active
- Identified by Process ID (PID)

# Multiprocessing: The (Traditional) Reality



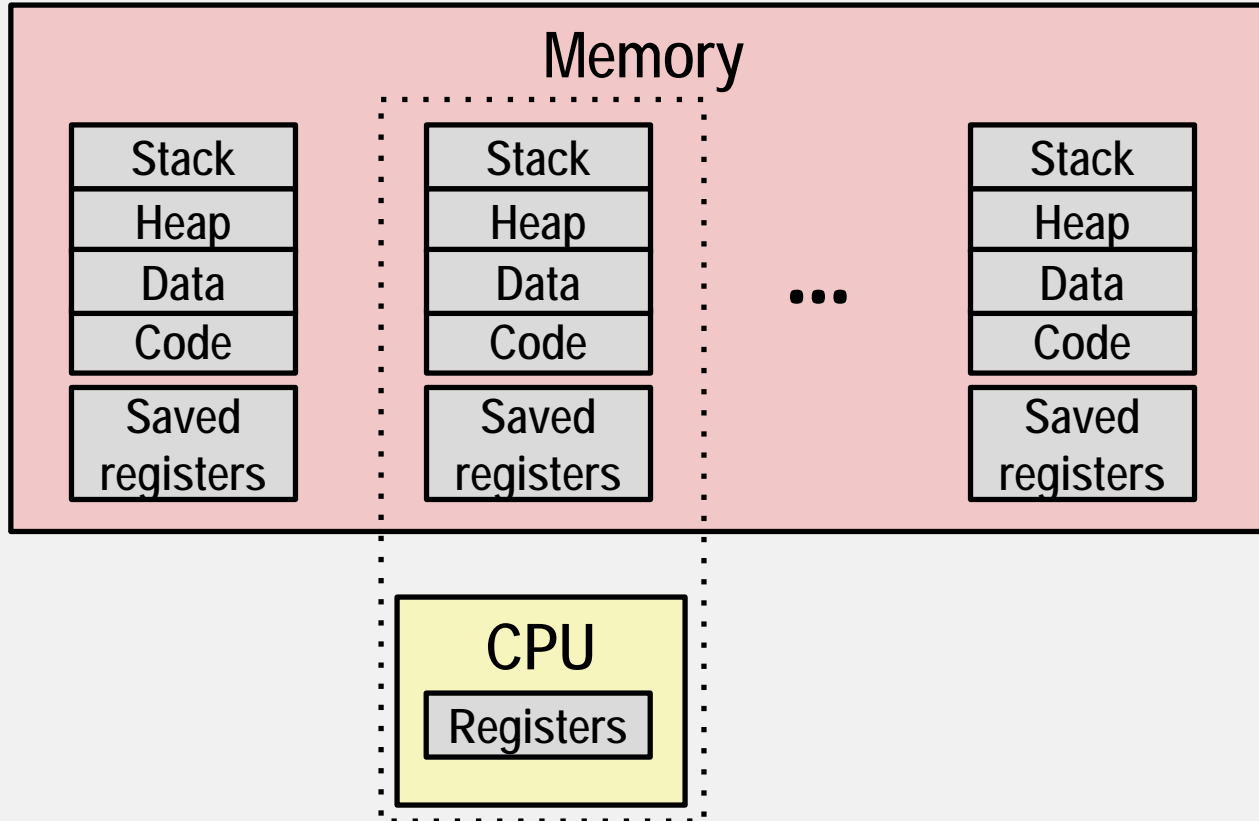
- **Single processor executes multiple processes concurrently**
  - Process executions interleaved (multitasking)
  - Address spaces managed by virtual memory system (later in course)
  - Register values for nonexecuting processes saved in memory

# Multiprocessing: The (Traditional) Reality



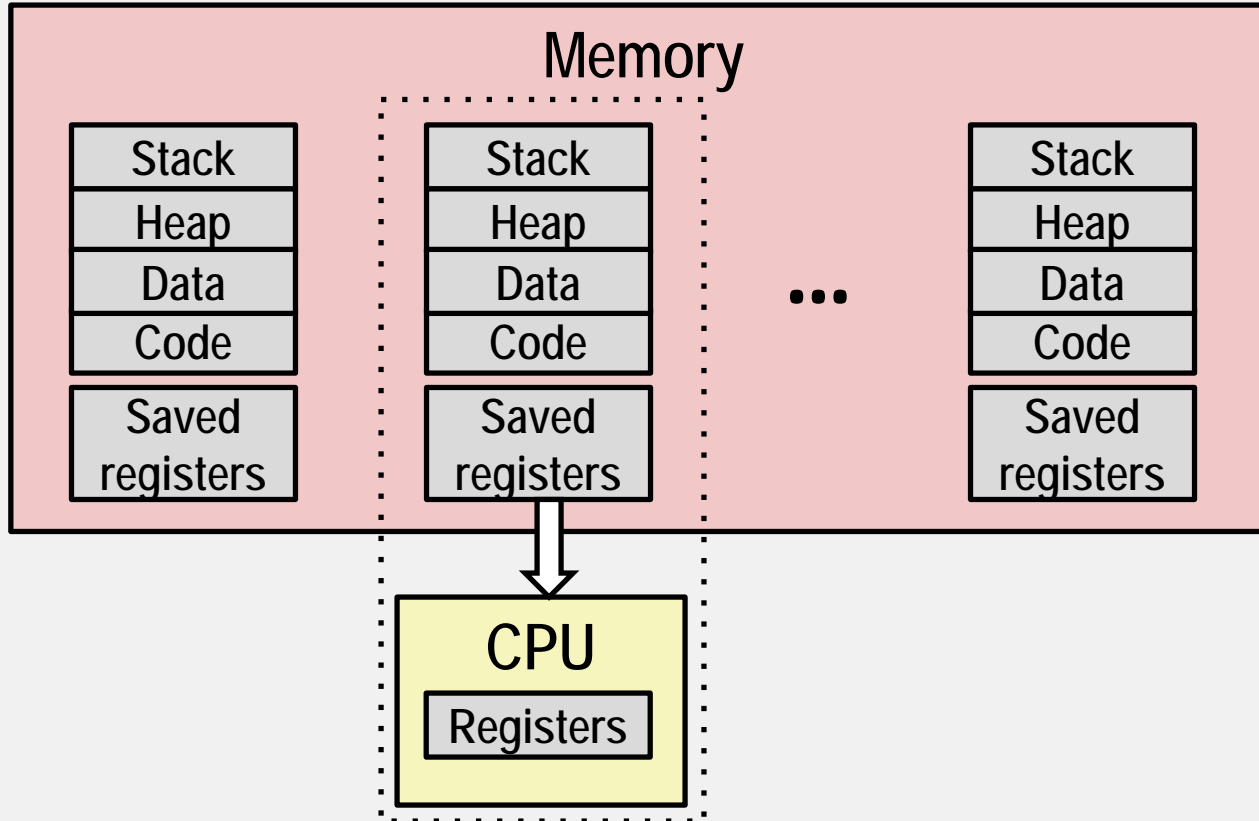
- Save current registers in memory

# Multiprocessing: The (Traditional) Reality



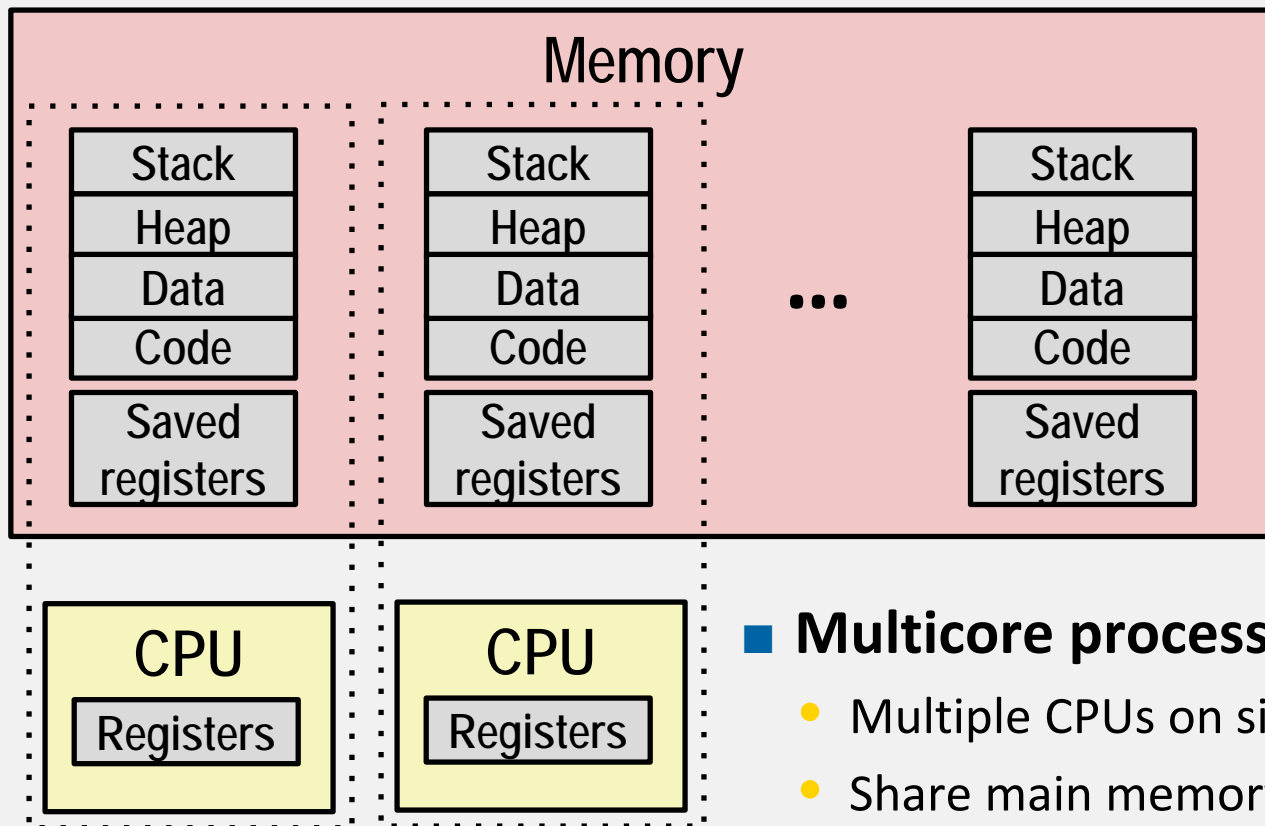
- Schedule next process for execution

# Multiprocessing: The (Traditional) Reality



- Load saved registers and switch address space (context switch)

# Multiprocessing: The (Modern) Reality



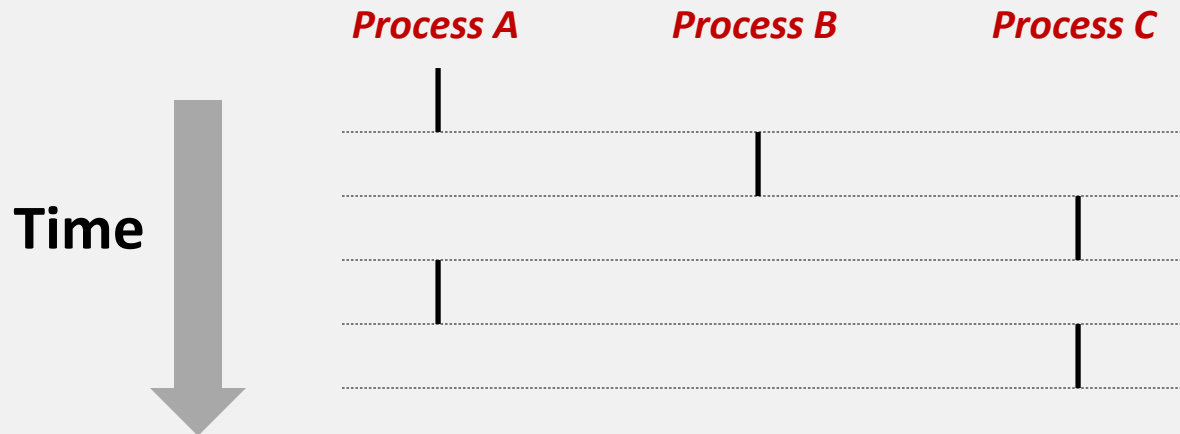
## ■ Multicore processors

- Multiple CPUs on single chip
- Share main memory (and some of the caches)
- Each can execute a separate process
  - Scheduling of processors onto cores done by kernel



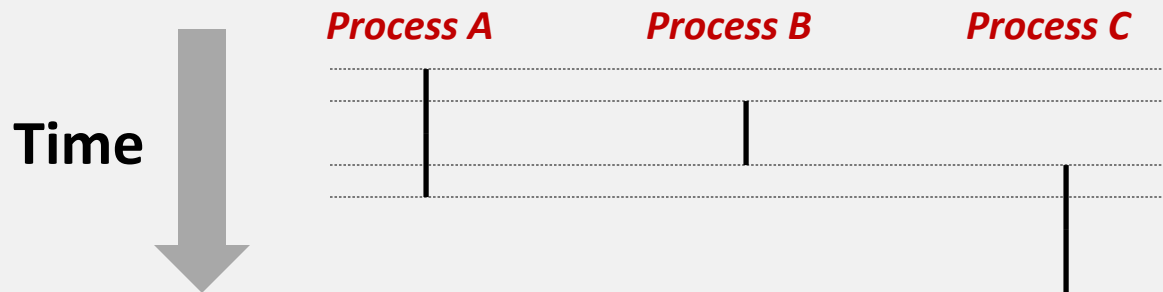
# Concurrent Processes

- Each process is a logical control flow.
- Two processes *run concurrently* (*are concurrent*) if their flows overlap in time
- Otherwise, they are *sequential*
- Examples (running on single core):
  - Concurrent: A & B, A & C
  - Sequential: B & C



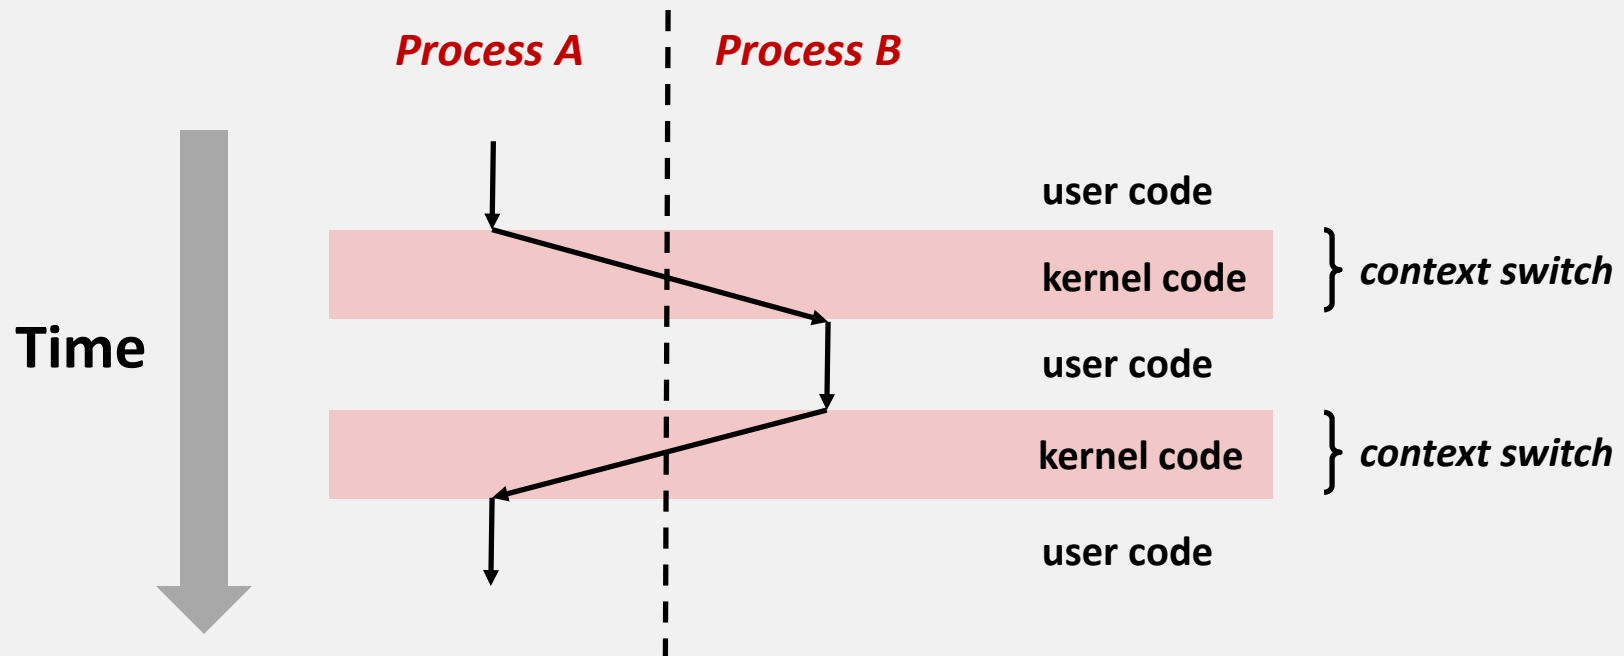
# User View of Concurrent Processes

- Control flows for concurrent processes are physically disjoint in time
- However, we can think of concurrent processes as running in parallel with each other



# Context Switching

- Processes are managed by a shared chunk of memory-resident OS code called the *kernel*
  - Important: the kernel is not a separate process, but rather runs as part of some existing process.
- Control flow passes from one process to another via a *context switch*



# Context Switching

## ■ Examples of Causes

- OS scheduling processes
- Kernel executing a syscall for user
  - If syscall blocks (read from disk), then let another process run
- Interrupts